

```
1 !pip install yfinance pandas matplotlib scikit-learn
```

```

Requirement already satisfied: yfinance in /usr/local/lib/python3.10/dist-packages (0.2.43)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.1.4)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.3.2)
Requirement already satisfied: numpy>=1.16.5 in /usr/local/lib/python3.10/dist-packages (from yfinance) (1.26.4)
Requirement already satisfied: requests>=2.31 in /usr/local/lib/python3.10/dist-packages (from yfinance) (2.32.3)
Requirement already satisfied: multitasking>=0.0.7 in /usr/local/lib/python3.10/dist-packages (from yfinance) (0.0.11)
Requirement already satisfied: lxml>=4.9.1 in /usr/local/lib/python3.10/dist-packages (from yfinance) (4.9.4)
Requirement already satisfied: platformdirs>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from yfinance) (4.3.2)
Requirement already satisfied: pytz>=2022.5 in /usr/local/lib/python3.10/dist-packages (from yfinance) (2024.1)
Requirement already satisfied: frozendict>=2.3.4 in /usr/local/lib/python3.10/dist-packages (from yfinance) (2.4.4)
Requirement already satisfied: peewee>=3.16.2 in /usr/local/lib/python3.10/dist-packages (from yfinance) (3.17.6)
Requirement already satisfied: beautifulsoup4>=4.11.1 in /usr/local/lib/python3.10/dist-packages (from yfinance) (4.12.3)
Requirement already satisfied: html5lib>=1.1 in /usr/local/lib/python3.10/dist-packages (from yfinance) (1.1)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2024.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.3.0)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.53.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.7)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (24.1)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (3.1.4)
Requirement already satisfied: scipy>=1.5.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.13.1)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.5.0)
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4>=4.11.1->yfinance) (2.6)
Requirement already satisfied: six>=1.9 in /usr/local/lib/python3.10/dist-packages (from html5lib>=1.1->yfinance) (1.16.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from html5lib>=1.1->yfinance) (0.5.1)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.31->yfinance) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.31->yfinance) (3.8)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.31->yfinance) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.31->yfinance) (2024.8.30)

```

```

1 #Import Libraries
2 import yfinance as yf
3 import pandas as pd
4 import matplotlib.pyplot as plt

```

 Generate

a slider using jupyter widgets





Close

```


1 #Download S&P500 Data
2 sp500 = yf.download('^GSPC', start='2000-01-01', end='2023-09-01')
3 sp500.head()

```

 [*****100%*****] 1 of 1 completed

	Open	High	Low	Close	Adj Close	Volume	
Date							
2000-01-03	1469.250000	1478.000000	1438.359985	1455.219971	1455.219971	931800000	
2000-01-04	1455.219971	1455.219971	1397.430054	1399.420044	1399.420044	1009000000	

```
1 # Download Nasdaq and Dow Jones data
2 nasdaq = yf.download('^IXIC', start='2000-01-01', end='2023-09-01')
3 dow_jones = yf.download('^DJI', start='2000-01-01', end='2023-09-01')
```

 [*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed

Next steps: | [Generate code with sn500](#) | | [View recommended plots](#) | | [New interactive sheet](#) |

```
1 #Download Apple and Google Stock Data
2 apple = yf.download('AAPL', start='2000-01-01', end='2023-09-01')
3 google = yf.download('GOOG', start='2000-01-01', end='2023-09-01')
4
```

 [*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed

```
1 sp500_close = sp500[['Adj Close']].rename(columns={'Adj Close': 'SP500'})
2 nasdaq_close = nasdaq[['Adj Close']].rename(columns={'Adj Close': 'Nasdaq'})
3 apple_close = apple[['Adj Close']].rename(columns={'Adj Close': 'Apple'})
```

```
1 #Merge Datasets
2 merged_data = pd.concat([sp500_close, nasdaq_close, apple_close], axis=1).dropna()
```

 Generate

create a dataframe with 2 columns and 10 rows



Close

```
1 #Plot the Data to Visually Compare Stock Performance
2 merged_data.plot(figsize=(10, 6))
3 plt.title('Stock Price Comparison')
4 plt.show()
5
```



Stock Price Comparison



```
1 #Daily Returns
```

```
2 merged_data['SP500_Return'] = merged_data['SP500'].pct_change()
```

```
3 merged_data['Nasdaq_Return'] = merged_data['Nasdaq'].pct_change()
```

```
4 merged_data['Apple_Return'] = merged_data['Apple'].pct_change()
```

```
5
```

```
1 # Weekly returns (resample data to weekly frequency)
```

```
2 merged_data['SP500_Weekly_Return'] = merged_data['SP500'].resample('W').ffill().pct_change()
```

```
3 merged_data['Nasdaq_Weekly_Return'] = merged_data['Nasdaq'].resample('W').ffill().pct_change()
```

```
4 merged_data['Apple_Weekly_Return'] = merged_data['Apple'].resample('W').ffill().pct_change()
```

```
5
```

```
1 # 50-day and 200-day moving averages for SP500
```

```
2 merged_data['SP500_50_MA'] = merged_data['SP500'].rolling(window=50).mean()
```

```
3 merged_data['SP500_200_MA'] = merged_data['SP500'].rolling(window=200).mean()
```

```
4
```

```
5 # 50-day and 200-day moving averages for Nasdaq and Apple
```

```
6 merged_data['Nasdaq_50_MA'] = merged_data['Nasdaq'].rolling(window=50).mean()
```

```
7 merged_data['Nasdaq_200_MA'] = merged_data['Nasdaq'].rolling(window=200).mean()
```

```
8 merged_data['Apple_50_MA'] = merged_data['Apple'].rolling(window=50).mean()
```

```
9 merged_data['Apple_200_MA'] = merged_data['Apple'].rolling(window=200).mean()
```

```
10
```

```
1 # 30-day rolling volatility for SP500
```

```
2 merged_data['SP500_Volatility'] = merged_data['SP500_Return'].rolling(window=30).std()
```

```
3
```

```
4 # 30-day rolling volatility for Nasdaq and Apple
```

```
5 merged_data['Nasdaq_Volatility'] = merged_data['Nasdaq_Return'].rolling(window=30).std()
```

```
6 merged_data['Apple_Volatility'] = merged_data['Apple_Return'].rolling(window=30).std()
```

```
7
```

```
1 #Creating Training and Test Sets
```

```
2 train_data = merged_data[:'2019']
```

```
3 test_data = merged_data['2020:']
```

```
1 from sklearn.linear_model import LinearRegression
```

```
2 # Combine features and target into one DataFrame for alignment
```

```
3 train_data_combined = pd.concat([train_data[['Nasdaq_Return', 'Apple_Return', 'Nasdaq_Volatility', 'Apple_Volatility']], train_data['SP500_Return']], axis=1).dropna()
```

```
4
```

```

5 # Separate back into features and target
6 X_train_clean = train_data_combined[['Nasdaq_Return', 'Apple_Return', 'Nasdaq_Volatility', 'Apple_Volatility']]
7 y_train_clean = train_data_combined['SP500_Return']
8
9 # Train the model on the cleaned data
10 model = LinearRegression()
11 model.fit(X_train_clean, y_train_clean)
12

```




LinearRegression
LinearRegression()

```

1 # Testing the model
2 X_test = test_data[['Nasdaq_Return', 'Apple_Return', 'Nasdaq_Volatility', 'Apple_Volatility']].dropna()
3 y_test = test_data['SP500_Return'].dropna()
4
5 # Predict SP500 returns
6 y_pred = model.predict(X_test)
7
8 # Set the plot size
9 plt.figure(figsize=(12, 6))
10 plt.plot(y_test.index, y_test, label='Actual SP500 Returns', alpha=0.75, color='blue')
11 plt.plot(y_test.index, y_pred, label='Predicted SP500 Returns', linestyle='--', alpha=0.75, color='orange')
12
13 plt.title('Actual vs Predicted SP500 Returns (2020-2023)', fontsize=16)
14 plt.xlabel('Date', fontsize=12)
15 plt.ylabel('Returns', fontsize=12)
16
17 plt.grid(True, linestyle='--', alpha=0.6)
18 plt.legend()
19 plt.show()
20

```



Predictive Analysis on S&P 500 Returns (2020-2023)

The model does a pretty good job of capturing the overall trend in S&P 500 returns. You can see this by looking at how closely the actual returns (blue line) and the predicted returns (orange dashed line) follow each other over time. This suggests that the model is able to pick up on general market movement patterns, especially between mid-2020 and 2023, where the predicted returns track the actual returns quite well.

Performance During High Volatility:

However, the model struggles a bit during periods of high market volatility. This is especially noticeable during the early 2020 market crash caused by the COVID-19 pandemic. While the model does manage to follow the overall direction of returns, it tends to underestimate just how sharp the drops and rebounds are. This suggests that the model might need additional features or a more complex structure to better handle sudden market shifts.

Post-Volatility Stabilization:

After the initial volatility in early 2020, the model's performance improves a lot. The predicted returns start to closely match the actual returns, showing that the model is able to adapt and perform better when the market is more stable. From 2021 to 2023, the alignment between actual and predicted returns becomes much tighter, indicating greater accuracy in predicting returns when there aren't as many extreme fluctuations.

Impact of Volatility and Returns Features:

Including volatility as a feature seems to have helped stabilize the model's predictions. Even during periods of higher volatility, like in 2022, the model does a decent job of tracking actual returns, though it still tends to underestimate the extreme swings, both upward and downward.

Overall, the model does well in capturing the general trend of S&P 500 returns and is pretty accurate in stable market conditions. However, there's room for improvement, especially when it comes to modeling extreme market events.

```
1 #Volume Data
2 merged_data['SP500_Volume'] = sp500['Volume']
3 merged_data['Nasdaq_Volume'] = nasdaq['Volume']
```

```
1 !pip install pandas_ta
2
```

```
Requirement already satisfied: pandas_ta in /usr/local/lib/python3.10/dist-packages (0.3.14b0)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from pandas_ta) (2.1.4)
Requirement already satisfied: numpy<2, >=1.22.4 in /usr/local/lib/python3.10/dist-packages (from pandas->pandas_ta) (1.26.4)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->pandas_ta) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->pandas_ta) (2024.1)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->pandas_ta) (2024.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->pandas_ta) (1.16.0)
```

```
1 import pandas_ta as ta
2
3 # Calculate RSI for SP500, Nasdaq, and Apple
4 merged_data['SP500_RSI'] = ta.rsi(merged_data['SP500'], length=14)
5 merged_data['Nasdaq_RSI'] = ta.rsi(merged_data['Nasdaq'], length=14)
6 merged_data['Apple_RSI'] = ta.rsi(merged_data['Apple'], length=14)
7
```

```
1 merged_data['SP500_Volume'] = sp500['Volume']
2 merged_data['Nasdaq_Volume'] = nasdaq['Volume']
3 merged_data['Apple_Volume'] = apple['Volume']
4
```

```
1 !pip install fredapi
2
```

```
Requirement already satisfied: fredapi in /usr/local/lib/python3.10/dist-packages (0.5.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from fredapi) (2.1.4)
Requirement already satisfied: numpy<2, >=1.22.4 in /usr/local/lib/python3.10/dist-packages (from pandas->fredapi) (1.26.4)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->fredapi) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->fredapi) (2024.1)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->fredapi) (2024.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->fredapi) (1.16.0)
```

```
1 # Fetching Economic Data such as Interest Rates, Inflation Rates, or GDP Growth Rates from the FRED API
2 from fredapi import Fred
3 fred = Fred(api_key='ed1116aca0b73c669e151ebf098b86d1')
4 # Fetch daily interest rate data
5 interest_rate = fred.get_series('DFF', start_date='2000-01-01', end_date='2023-09-01')
6
7 # Convert the series to a DataFrame and assign a column name
8 interest_rate_df = interest_rate.to_frame(name='Interest_Rate')
9
10 # Merge the interest rate data with the merged_data
11 merged_data = merged_data.merge(interest_rate_df, how='left', left_index=True, right_index=True)
```

12
13
14

```
1 print(merged_data.isnull().sum())
2
```

```
SP500      0
Nasdaq     0
Apple      0
SP500_Return  1
Nasdaq_Return  1
Apple_Return  1
SP500_Weekly_Return  5954
Nasdaq_Weekly_Return  5954
Apple_Weekly_Return  5954
SP500_50_MA  49
SP500_200_MA  199
Nasdaq_50_MA  49
Nasdaq_200_MA  199
Apple_50_MA  49
Apple_200_MA  199
SP500_Volatility  30
Nasdaq_Volatility  30
Apple_Volatility  30
SP500_Volume  0
Nasdaq_Volume  0
SP500_RSI  14
Nasdaq_RSI  14
Apple_RSI  14
Apple_Volume  0
Interest_Rate  0
dtype: int64
```

```
1 # Fill the single missing value in returns
2 merged_data['SP500_Return'].fillna(method='ffill', inplace=True)
3 merged_data['Nasdaq_Return'].fillna(method='ffill', inplace=True)
4 merged_data['Apple_Return'].fillna(method='ffill', inplace=True)
5
6 # Drop weekly returns if they are not needed
7 merged_data.drop(columns=['SP500_Weekly_Return', 'Nasdaq_Weekly_Return', 'Apple_Weekly_Return'], inplace=True)
8
9 # Fill missing values in moving averages by forward filling
10 merged_data['SP500_50_MA'].fillna(method='ffill', inplace=True)
11 merged_data['SP500_200_MA'].fillna(method='ffill', inplace=True)
12 merged_data['Nasdaq_50_MA'].fillna(method='ffill', inplace=True)
13 merged_data['Nasdaq_200_MA'].fillna(method='ffill', inplace=True)
14 merged_data['Apple_50_MA'].fillna(method='ffill', inplace=True)
15 merged_data['Apple_200_MA'].fillna(method='ffill', inplace=True)
16
17 # Alternatively, if you're okay with dropping initial rows for MA calculations:
18 merged_data = merged_data.dropna(subset=['SP500_200_MA', 'Nasdaq_200_MA', 'Apple_200_MA'])
19
20 # Forward-fill missing volatility values
21 merged_data['SP500_Volatility'].fillna(method='ffill', inplace=True)
22 merged_data['Nasdaq_Volatility'].fillna(method='ffill', inplace=True)
23 merged_data['Apple_Volatility'].fillna(method='ffill', inplace=True)
```

```

24
25 merged_data['SP500_RSI'].fillna(method='ffill', inplace=True)
26 merged_data['Nasdaq_RSI'].fillna(method='ffill', inplace=True)
27 merged_data['Apple_RSI'].fillna(method='ffill', inplace=True)
28
29 # Fill missing values for economic data like interest rate with zero
30 merged_data['Interest_Rate'].fillna(0, inplace=True)
31
32 # Check for missing values after filling
33 print(merged_data.isnull().sum())
34
35

```

```

SP500      0
Nasdaq     0
Apple      0
SP500_Return  1
Nasdaq_Return  1
Apple_Return  1
SP500_50_MA  49
SP500_200_MA 199
Nasdaq_50_MA  49
Nasdaq_200_MA 199
Apple_50_MA   49
Apple_200_MA  199
SP500_Volatility  30
Nasdaq_Volatility  30
Apple_Volatility  30
SP500_Volume    0
Nasdaq_Volume    0
SP500_RSI       14
Nasdaq_RSI       14
Apple_RSI        14
Apple_Volume     0
Interest_Rate    0
dtype: int64
<ipython-input-149-d09895129129>:2: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['SP500_Return'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:3: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['Nasdaq_Return'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:4: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['Apple_Return'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:10: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['SP500_50_MA'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:11: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['SP500_200_MA'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:12: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['Nasdaq_50_MA'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:13: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['Nasdaq_200_MA'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:14: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['Apple_50_MA'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:15: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['Apple_200_MA'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:21: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['SP500_Volatility'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:22: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['Nasdaq_Volatility'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:23: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['Apple_Volatility'].fillna(method='ffill', inplace=True)

```



```

<ipython-input-149-d09895129129>:25: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['SP500_RSI'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:26: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['Nasdaq_RSI'].fillna(method='ffill', inplace=True)
<ipython-input-149-d09895129129>:27: FutureWarning: Series.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.
merged_data['Apple_RSI'].fillna(method='ffill', inplace=True)

```

```

1 #Drop all Null Values
2 merged_data = merged_data.dropna()
3

```

```

1 print(train_data.columns)

```

```

Index(['SP500', 'Nasdaq', 'Apple', 'SP500_Return', 'Nasdaq_Return',
      'Apple_Return', 'SP500_Weekly_Return', 'Nasdaq_Weekly_Return',
      'Apple_Weekly_Return', 'SP500_50_MA', 'SP500_200_MA', 'Nasdaq_50_MA',
      'Nasdaq_200_MA', 'Apple_50_MA', 'Apple_200_MA', 'SP500_Volatility',
      'Nasdaq_Volatility', 'Apple_Volatility'],
      dtype='object')

```

```

1 train_data = merged_data[:'2019']
2 test_data = merged_data['2020:']
3
4 # Include new features like RSI, volume, and interest rates in the training data
5 features = ['Nasdaq_Return', 'Apple_Return', 'Nasdaq_Volatility', 'Apple_Volatility',
6            'SP500_RSI', 'Nasdaq_RSI', 'Apple_RSI', 'SP500_Volume', 'Nasdaq_Volume',
7            'Apple_Volume', 'Interest_Rate']
8
9 X_train_clean = train_data[features]
10 y_train_clean = train_data['SP500_Return']
11
12 X_test_clean = test_data[features]
13 y_test_clean = test_data['SP500_Return']
14

```

```

1 #Random Forrest
2 from sklearn.ensemble import RandomForestRegressor
3
4 rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
5 rf_model.fit(X_train_clean, y_train_clean)
6 y_rf_pred = rf_model.predict(X_test_clean)
7

```

```

1 #XG Boost
2 import xgboost as xgb
3
4 xg_model = xgb.XGBRegressor(objective='reg:squarederror', n_estimators=100)
5 xg_model.fit(X_train_clean, y_train_clean)
6 y_xg_pred = xg_model.predict(X_test_clean)
7

```

```

1 #LSTM
2 from keras.models import Sequential

```

```
3 from keras.layers import LSTM, Dense
4
5 # Reshape data for LSTM
6 X_train_lstm = X_train_clean.values.reshape((X_train_clean.shape[0], 1, X_train_clean.shape[1]))
7 X_test_lstm = X_test_clean.values.reshape((X_test_clean.shape[0], 1, X_test_clean.shape[1]))
8
9 model = Sequential()
10 model.add(LSTM(50, activation='relu', input_shape=(X_train_lstm.shape[1], X_train_lstm.shape[2])))
11 model.add(Dense(1))
12 model.compile(optimizer='adam', loss='mse')
13
14 model.fit(X_train_lstm, y_train_clean, epochs=100, batch_size=32)
15 y_lstm_pred = model.predict(X_test_lstm)
16
```



```

Epoch 96/100
151/151 ————— 0s 2ms/step - loss: 10833733632.0000
Epoch 97/100
151/151 ————— 0s 2ms/step - loss: 87446839296.0000
Epoch 98/100
151/151 ————— 1s 2ms/step - loss: 31228592128.0000
Epoch 99/100
151/151 ————— 0s 2ms/step - loss: 180687110144.0000
Epoch 100/100
151/151 ————— 1s 3ms/step - loss: 19393269760.0000
29/29 ————— 1s 16ms/step

```

```

1 from sklearn.metrics import mean_absolute_error, mean_squared_error
2 import numpy as np
3
4 # Random Forest Evaluation
5 mae_rf = mean_absolute_error(y_test_clean, y_rf_pred)
6 rmse_rf = np.sqrt(mean_squared_error(y_test_clean, y_rf_pred))
7
8 # XGBoost Evaluation
9 mae_xg = mean_absolute_error(y_test_clean, y_xg_pred)
10 rmse_xg = np.sqrt(mean_squared_error(y_test_clean, y_xg_pred))
11
12 # LSTM Evaluation
13 mae_lstm = mean_absolute_error(y_test_clean, y_lstm_pred)
14 rmse_lstm = np.sqrt(mean_squared_error(y_test_clean, y_lstm_pred))
15
16 print(f'Random Forest MAE: {mae_rf}, RMSE: {rmse_rf}')
17 print(f'XGBoost MAE: {mae_xg}, RMSE: {rmse_xg}')
18 print(f'LSTM MAE: {mae_lstm}, RMSE: {rmse_lstm}')
19

```

```

→ Random Forest MAE: 0.0035508102184967534, RMSE: 0.005024369435059261
XGBoost MAE: 0.0037702658392709033, RMSE: 0.005588085697649882
LSTM MAE: 2381955.905770924, RMSE: 5082449.148507779

```

✓ Model Performance Explanation

After training three models—**Random Forest**, **XGBoost**, and **LSTM**—on the dataset, the performance metrics of each model were evaluated using **Mean Absolute Error (MAE)** and **Root Mean Squared Error (RMSE)**. The results are as follows:

- **Random Forest:**

- MAE: 0.00355
- RMSE: 0.00502

Random Forest performed the best out of the three models. The low MAE and RMSE values indicate that the model made relatively accurate predictions, with only small errors on average. This suggests that Random Forest was able to capture the underlying patterns in the data well.

- **XGBoost:**

- MAE: 0.00377

- RMSE: 0.00559

XGBoost also performed well but slightly worse than Random Forest. The MAE and RMSE are a little higher, meaning that the model's predictions were a bit less accurate. However, it still managed to provide reasonably good predictions overall.

- **LSTM (Long Short-Term Memory):**


- MAE: 1,417,528.16
- RMSE: 3,220,159.74

The LSTM model performed very poorly compared to Random Forest and XGBoost. The extremely high MAE and RMSE suggest that the model struggled to learn meaningful patterns in the data. This could be due to several factors, such as improper data preprocessing, insufficient training, or an inappropriate model configuration. Further tuning and data handling would be needed to make the LSTM model perform better.

Overall, **Random Forest** emerged as the best-performing model for this task, while the **LSTM** model failed to provide meaningful predictions, likely due to issues with the setup.

```
1 test_data['Predicted_Return'] = y_rf_pred # or use y_xg_pred, y_lstm_pred
2 test_data['Signal'] = np.where(test_data['Predicted_Return'] > 0, 1, -1)
3
4 # Calculate strategy returns based on signals
5 test_data['Strategy_Return'] = test_data['Signal'] * test_data['SP500_Return']
6 cumulative_strategy_return = (test_data['Strategy_Return'] + 1).cumprod()
7
8 # Plot the strategy graph
9 plt.figure(figsize=(10, 6))
10 plt.plot(cumulative_strategy_return, label='Strategy Return')
11 plt.title('Backtested Strategy Return')
12 plt.legend()
13 plt.show()
14
```

```
<ipython-input-157-b2a70059db8b>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead  
  
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
test_data['Predicted_Return'] = y_rf_pred # or use y_xg_pred, y_lstm_pred  
<ipython-input-157-b2a70059db8b>:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead  
  
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
test_data['Signal'] = np.where(test_data['Predicted_Return'] > 0, 1, -1)  
<ipython-input-157-b2a70059db8b>:5: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead  
  
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
test_data['Strategy_Return'] = test_data['Signal'] * test_data['SP500_Return']
```



Backtested Strategy Performance Analysis:

The backtest of the strategy, based on predictions from the machine learning model, shows impressive exponential returns starting from mid-2021, with a sharp increase in profits extending into 2023. The slow and steady growth before mid-2021 reflects the strategy's lackluster performance during the earlier market conditions, possibly due to the model's inability to capture the market's volatile behavior during the COVID-19 pandemic. However, after 2021, the strategy shows strong compounding returns, indicating that the model was able to capture market trends effectively during this period.

Despite the substantial gains, the exponential growth in returns may suggest potential overfitting to historical data, and thus, further testing on real-time or out-of-sample data is recommended to assess the strategy's long-term robustness and performance in various market conditions.

```

1 # Fetch the latest data for each asset separately for clarity
2 sp500_data = yf.download('^GSPC', start='2023-09-01', end='2024-01-01')
3 nasdaq_data = yf.download('^IXIC', start='2023-09-01', end='2024-01-01')
4 apple_data = yf.download('AAPL', start='2023-09-01', end='2024-01-01')
5
6 # Reset the index for each DataFrame to get a flat index
7 sp500_data = sp500_data.reset_index()
8 nasdaq_data = nasdaq_data.reset_index()
9 apple_data = apple_data.reset_index()
10

```

```

[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completed

```

```

1 # 1. Calculate daily returns for each asset
2 sp500_data['SP500_Return'] = sp500_data['Adj Close'].pct_change()
3 nasdaq_data['Nasdaq_Return'] = nasdaq_data['Adj Close'].pct_change()
4 apple_data['Apple_Return'] = apple_data['Adj Close'].pct_change()
5
6 # 2. Calculate 30-day rolling volatility for Nasdaq and Apple
7 nasdaq_data['Nasdaq_Volatility'] = nasdaq_data['Nasdaq_Return'].rolling(window=30).std()
8 apple_data['Apple_Volatility'] = apple_data['Apple_Return'].rolling(window=30).std()
9
10 # 3. Calculate RSI using pandas_ta for each asset
11 sp500_data['SP500_RSI'] = ta.rsi(sp500_data['Adj Close'], length=14)
12 nasdaq_data['Nasdaq_RSI'] = ta.rsi(nasdaq_data['Adj Close'], length=14)
13 apple_data['Apple_RSI'] = ta.rsi(apple_data['Adj Close'], length=14)
14
15 # 4. Add volume data
16 sp500_data['SP500_Volume'] = sp500_data['Volume']
17 nasdaq_data['Nasdaq_Volume'] = nasdaq_data['Volume']
18 apple_data['Apple_Volume'] = apple_data['Volume']
19
20 # 5. Keep only the Date and the required columns for each DataFrame
21 sp500_features = sp500_data[['Date', 'SP500_Return', 'SP500_RSI', 'SP500_Volume']]
22 nasdaq_features = nasdaq_data[['Date', 'Nasdaq_Return', 'Nasdaq_Volatility', 'Nasdaq_RSI', 'Nasdaq_Volume']]
23 apple_features = apple_data[['Date', 'Apple_Return', 'Apple_Volatility', 'Apple_RSI', 'Apple_Volume']]
24
25 # Print to check the data structure
26 print(sp500_features.head())
27 print(nasdaq_features.head())
28 print(apple_features.head())
29

```

```

[*****]

```

	Date	SP500_Return	SP500_RSI	SP500_Volume
0	2023-09-01	NaN	NaN	3246260000
1	2023-09-05	-0.004194	NaN	3526250000
2	2023-09-06	-0.006972	NaN	3418850000
3	2023-09-07	-0.003211	NaN	3763760000
4	2023-09-08	0.001427	NaN	3259290000

```

Date Nasdaq_Return Nasdaq_Volatility Nasdaq_RSI Nasdaq_Volume

```


0	2023-09-01	NaN	NaN	NaN	4033960000
1	2023-09-05	-0.000774	NaN	NaN	4379790000
2	2023-09-06	-0.010590	NaN	NaN	4215320000
3	2023-09-07	-0.008913	NaN	NaN	4320830000
4	2023-09-08	0.000924	NaN	NaN	4160360000

	Date	Apple_Return	Apple_Volatility	Apple_RSI	Apple_Volume
0	2023-09-01	NaN	NaN	NaN	45732600
1	2023-09-05	0.001267	NaN	NaN	45280000
2	2023-09-06	-0.035793	NaN	NaN	81755800
3	2023-09-07	-0.029249	NaN	NaN	112488800
4	2023-09-08	0.003492	NaN	NaN	65551300

```

1 # 6. Fetch interest rate data from FRED
2 fred = Fred(api_key='ed1116aca0b73c669e151ebf098b86d1')
3
4 # Fetch daily interest rate data
5 interest_rate = fred.get_series('DFF', start='2023-09-01', end='2024-01-01')
6
7 # Convert the interest rate series into a DataFrame and reset the index
8 interest_rate_df = pd.DataFrame(interest_rate, columns=['Interest_Rate'])
9 interest_rate_df.index = pd.to_datetime(interest_rate_df.index)
10 interest_rate_df = interest_rate_df.reset_index()
11 interest_rate_df.columns = ['Date', 'Interest_Rate']
12
13 # Print to check interest rate data
14 print(interest_rate_df.head())
15

```




	Date	Interest_Rate
0	1954-07-01	1.13
1	1954-07-02	1.25
2	1954-07-03	1.25
3	1954-07-04	1.25
4	1954-07-05	0.88

```

1 # 7. Merge SP500, Nasdaq, Apple, and Interest Rate data
2 merged_data = sp500_features.merge(nasdaq_features, on='Date').merge(apple_features, on='Date').merge(interest_rate_df, on='Date')
3
4 # Drop any rows with missing values after merging
5 merged_data = merged_data.dropna()
6
7 # Print the merged data to check the final structure
8 print(merged_data.head())
9

```



	Date	SP500_Return	SP500_RSI	SP500_Volume	Nasdaq_Return	\
30	2023-10-16	0.010594	51.869604	3409960000	0.011990	
31	2023-10-17	-0.000098	51.800114	3794850000	-0.002523	
32	2023-10-18	-0.013400	43.283656	3686030000	-0.016215	
33	2023-10-19	-0.008483	38.973738	3969730000	-0.009623	
34	2023-10-20	-0.012585	33.663536	4004030000	-0.015347	

	Nasdaq_Volatility	Nasdaq_RSI	Nasdaq_Volume	Apple_Return	\
30	0.009990	52.287030	4308690000	-0.000727	
31	0.009993	50.973425	4417640000	-0.008785	
32	0.010227	43.438875	4617140000	-0.007395	
33	0.010246	39.745157	5014790000	-0.002161	

34	0.010548	34.722739	4622840000	-0.014704
	Apple_Volatility	Apple_RSI	Apple_Volume	Interest_Rate
30	0.013153	52.040961	52517000	5.33
31	0.013199	47.762466	57549400	5.33
32	0.011631	44.476756	54764400	5.33
33	0.010365	43.541121	59302900	5.33
34	0.010661	37.736299	64189300	5.33

```

1 # 8. Prepare the features for model prediction
2 features = ['Nasdaq_Return', 'Apple_Return', 'Nasdaq_Volatility', 'Apple_Volatility',
3             'SP500_RSI', 'Nasdaq_RSI', 'Apple_RSI', 'SP500_Volume', 'Nasdaq_Volume',
4             'Apple_Volume', 'Interest_Rate']
5
6 latest_features = merged_data[features].values
7
8 # Print to check the shape of the feature matrix
9 print(latest_features.shape)
10

```

→ (53, 11)

```

1 # Predict future returns with Random Forest model
2 predicted_rf_returns = rf_model.predict(latest_features)
3
4 # For XGBoost:
5 predicted_xg_returns = xg_model.predict(latest_features)
6
7 # For LSTM (reshape the data before prediction):
8 latest_features_reshaped = latest_features.reshape((latest_features.shape[0], 1, latest_features.shape[1]))
9 predicted_lstm_returns = model.predict(latest_features_reshaped)
10

```

→ 2/2 ————— 0s 8ms/step

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:465: UserWarning: X does not have valid feature names, but RandomForestRegressor was fitted with feature names
warnings.warn(

```

1 # Add predicted returns to the merged_data DataFrame
2 merged_data['Predicted>Returns'] = predicted_rf_returns
3
4 # You can also plot other model results like predicted_xg_returns or predicted_lstm_returns
5 # Just replace 'predicted_rf_returns' with the model-specific results
6
7 # Optional: If you want to compare with actual SP500 returns (for example), we can keep the SP500 actual returns
8 merged_data['Actual>Returns'] = merged_data['SP500_Return']
9
10 # Print to see the DataFrame structure before plotting
11 print(merged_data[['Date', 'Actual>Returns', 'Predicted>Returns']].head())
12

```

→

	Date	Actual>Returns	Predicted>Returns
30	2023-10-16	0.010594	0.008920
31	2023-10-17	-0.000098	-0.001218
32	2023-10-18	-0.013400	-0.009756
33	2023-10-19	-0.008483	-0.006299

34 2023-10-20

-0.012585

-0.011777

```
1 # Set the plot size
2 plt.figure(figsize=(12, 6))
3
4 # Plot actual returns (for example, SP500)
5 plt.plot(merged_data['Date'], merged_data['Actual_Returns'], label='Actual SP500 Returns', color='blue', alpha=0.75)
6
7 # Plot predicted returns (from Random Forest)
8 plt.plot(merged_data['Date'], merged_data['Predicted_Returns'], label='Predicted SP500 Returns (RF)', color='orange', linestyle='--', alpha=0.75)
9
10 # Add title and labels
11 plt.title('Actual vs Predicted SP500 Returns', fontsize=16)
12 plt.xlabel('Date', fontsize=12)
13 plt.ylabel('Returns', fontsize=12)
14
15 # Add grid for better readability
16 plt.grid(True, linestyle='--', alpha=0.6)
17
18 # Add legend
19 plt.legend()
20
21 # Show the plot
22 plt.show()
23
```

